

# From organically grown infrastructure, to a mature ML platform

OLIVIA STOICESCU, Sr. MLOps Engineer  
23 Feb 2023

# Agenda

---

- What is an ML Platform?
- ML Workflows
- From ML Workflows to a ML Platform
  - What is the Problem you are Solving?
  - Platform Interaction Modes
  - Platform Onboarding
  - Buy vs Build
  - The Path to Platform Enablement
- Conclusion

# WHAT is a Platform?

---

Team Topologies book introduces the concept of **Thinnest Viable Platform (TVP)**:

*A TVP is the smallest set of APIs, documentation, and tools needed to accelerate the teams developing modern software services and systems.*

# TVP

---

.... does this mean that ***anything*** can be considered a platform??



# What is a ML Platform?

---

**A platform which implements DevOps principles to ML workflows.**

...however, MLOps is not DevOps:

- ML is experimental in nature

- Testing a ML system is more involved

- Production models might decay due to changing data profiles

# ML Workflows

---

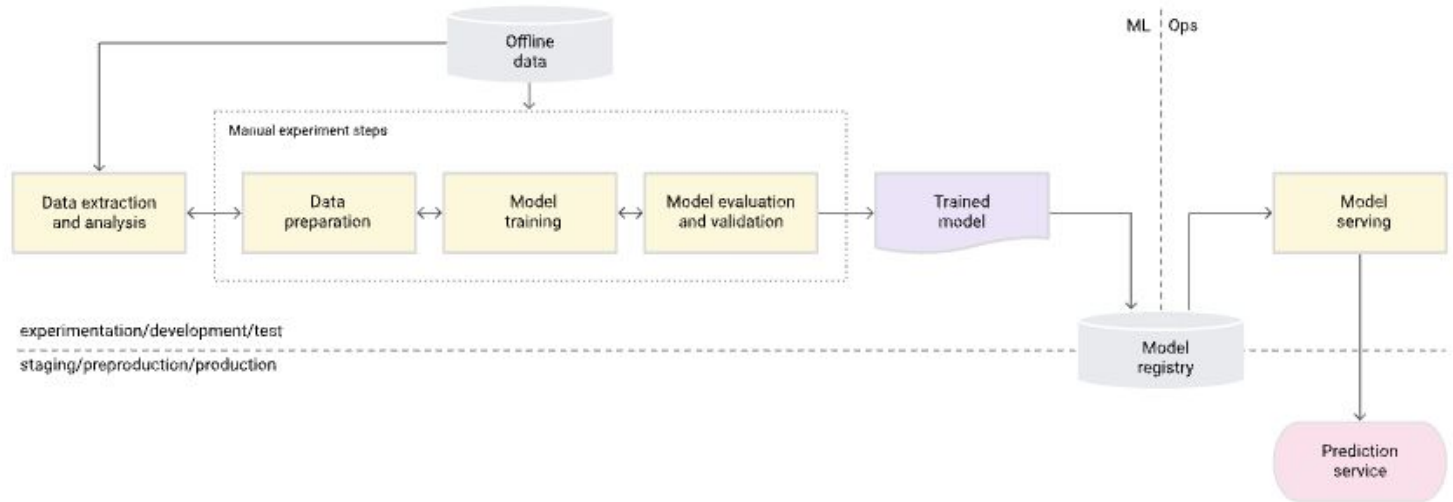
**ML workflows** include the following key assets: **code**, **models** and **data**.

... and the following stages: **train**, **test**, **package**, **deploy**, **monitor** and **maintain** models

# ML Workflow - no automation

1. The entire Machine Learning Pipeline is executed in the Experimentation Environment manually and on demand.
2. After the Model Artifact is created it is saved into a Model Registry
3. Model serving uses already present deployment procedures

Diagram source:  
<https://cloud.google.com/architecture/mlops-continuous-delivery-and-automation-pipelines-in-machine-learning>



# ML Workflows - automation

---

Implementing ML in a production environment doesn't only mean **deploying your model** as an **API** for prediction.

It means deploying an **ML pipeline** that can automate the retraining and deployment of new models.

Setting up a **CI/CD system** enables you to automatically test and deploy new pipeline implementations.

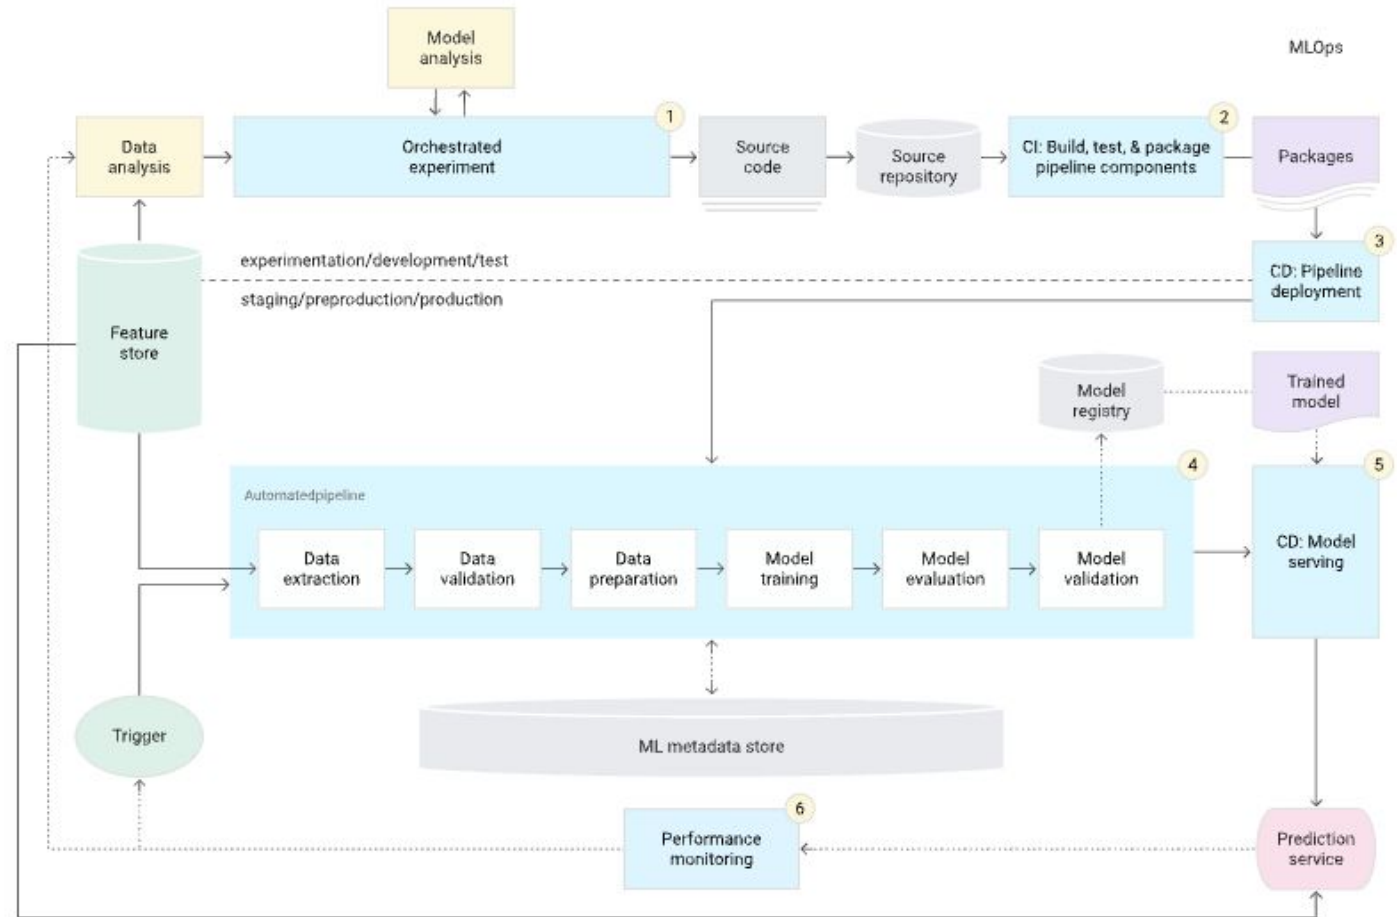
- This system lets you cope with rapid changes in your data and business environment.
- You can gradually implement these practices to help improve the automation of your ML system development and production.



# ML workflow - Mature

1. ML Pipelines are orchestrated but you can trigger them from the experimentation environment - Notebooks.
2. Continuous Integration Step for ML Pipelines
3. Continuous Delivery/ Deployment Step for ML Pipelines.
4. Automated triggering of the ML Pipeline.
5. ML Model Continuous Delivery.
6. Performance Monitoring.

Diagram source:  
<https://cloud.google.com/architecture/mlops-continuous-delivery-and-automation-pipelines-in-machine-learning>



# From ML Workflows to an ML Platform

---

# What problem are you solving?

---

👉 **Enable scaling** - allow the internal customers to become self-sufficient in their ML journey

👉 **Improving the ML output quality** - eg. moving from batch to real-time processing

👉 **Improve developer experience** - eg. automation, minimize cognitive load

# Platform Interaction Modes

---

Understanding who is the key **internal customer** will help understand the optimal interaction mode:

## **data-scientists, data analysts**

- can have have limited experience with the software life-cycle
- Lead the interaction via configuration

## **(ML) engineers**

- more experienced with software lifecycle and tooling
- Build apis and cli tools

# Platform Onboarding

---

The platform should aim to enable **fast onboarding**:

- Automate generating compliant projects (eg. via cookiecutter for python, archetypes for jvm)
- Standardize dependency management to scale fast to new projects (eg. via a parent pom for jvm, or using tools like poetry for python)
- Standardize CI/CD pipeline to scale fast to new projects (eg. via templates)
- Use gitOps to better support traceability and change management (eg. by using environment manifest repositories)

# Buy vs Build: Open-source

---



## Pros:

- Lots of free extensions, allowing to cover for missing functionality
- Flexibility in customization (you can decide which features to use or leave out)
- Community delivering case studies, tutorials, how-tos

## Cons:

- Slower adoption (installation and initial setup can be a bumpy road)
- Maintenance (eg. kubeflow relies on k8s as base platform)



# Buy vs Build: Managed

---

## Pros:

- Fast adoption
- Effort-free features
- Strategic partner (support)

## Cons:

- Vendor selection
- Vendor lock-in

# Buy vs Build

---

Most often, the discussion between managed and open-source comes down to **what is the most limited resource you have.**

- A managed MLOps platform minimizes the need for engineering resources but requires a certain amount of investment.
- While open source MLOps platforms are the other way around.

You can choose to apply a blend of the two and go for open-source on the components where in-house expertise is most aligned with the value streams




# The Path to Platform Enablement

Typically, we start from some initial ML development setup.

At some point, this setup outgrows its initial purpose and needs to evolve.

Diagram source:  
<https://www.youtube.com/watch?v=WaL3ZbLgMul>



# Final thoughts

---

Favour incremental changes

- Changes are frequent and so they are expected

Favour loose coupling

- build extensible abstraction layers, in order to freely move the technologies behind the scenes

Favour testability of your platform components by design

- embedded visibility - if this changes, what/where is the impact?
- lack of visibility hinders incremental changes

# Conclusion


---

Keep in mind your main objective for creating or evolving a ML platform

Prioritize building for the key internal customer first

Minimize adoption friction by building appropriate interfacing tools

What got you here won't get you there: skills and expertise that informed decisions at earlier stages, might not fit the current goals





```
if questions:
    try:
        answer()
    except RuntimeError:
        pass
else:
    print("Thank You.")
```